


Building deep learning training pipelines

Perumadura De Silva

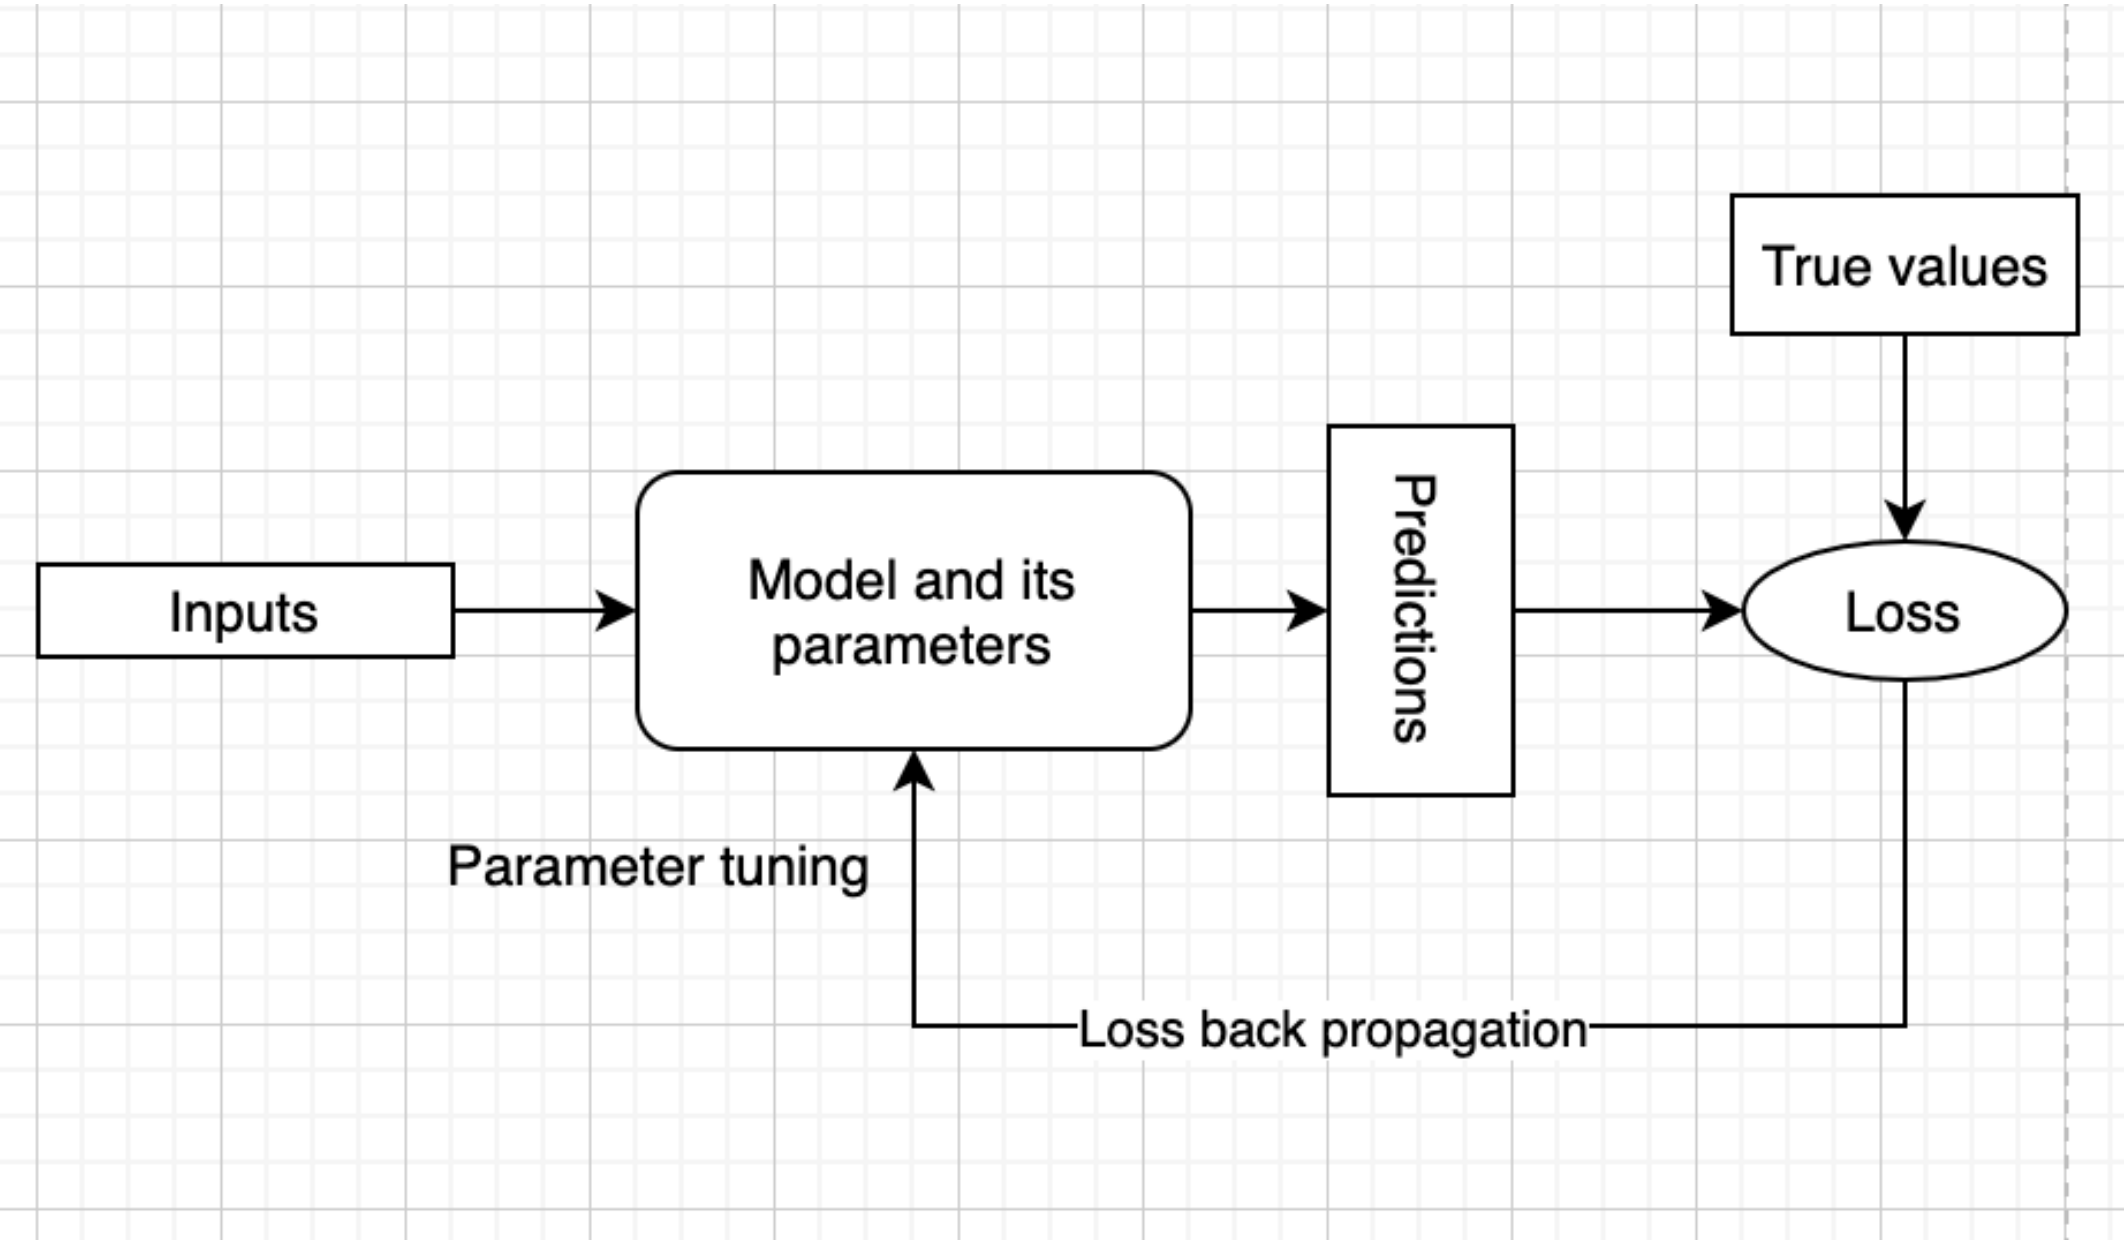
Content

- Structure of a deep learning training pipeline
- Training process of a neural network
- Introduction to model optimisation
- Deep neural network optimisation

Structure of a deep learning training pipeline



Ref.1 Multi-layer Neural Network



Components involved in supervise learning

Structure of a deep learning training pipeline

- Input layer:

- Inputs to the neural network such as images, audio, text embedding etc.

- Example:

Inputs
↓


0.5	0.2	0.1	0.33	0.51	0.14
0.32	0.53	0.34	0.34	0.13	0.44
0.14	0.56	0.44	0.41	0.46	0.65

← Number of data samples

- Each input sample is fed into the neural network:
 - Batch-wise (A set of samples together for a training iteration)
 - single example (one sample per training iteration)

Structure of a deep learning training pipeline

- Hidden layer:
 - Inputs connects to hidden layer through the weights
 - Extract the features of the inputs by means of weights and biases
 - This feature extraction encodes the information about an input
 - Then propagate this this information towards output layer
 - Activation function: Performs a nonlinear transformation on the inputs



Structure of a deep learning training pipeline

- Output layer:
 - Takes the inputs from the last hidden layer
 - Produce predictions
 - Similar to the inputs, the outputs has the same size in the batch dimension
 - The output dimension may differ

Training process of a neural network

- Three main components involved in training
 - Loss function
 - Compute a distance/encoding between the predictions and true values
 - Used in both supervised, semi supervised and self supervised cases
 - Example MEAN SQUARED ERROR:

$$mse = \frac{1}{2N} \sum_i^N (\hat{y} - y)^2$$

Training process of a neural network

- Neuron activation

- Introduces nonlinearity to the layer outputs

- Selection of the activation function is somewhat empirical

- Example (Sigmoid function): $\sigma(t) = \frac{1}{1 + e^{-t}}$


- Must be differentiable

- Optimisation strategy

- Two major components: Gradient descent and error back propagation

- Gradient descent: Update parameters so the loss descend towards a global minima

- Error backdrop: Propagate the loss gradients towards the parameters for gradient descent computation



Ref.4: Sigmoid function

Colab Exercise

Training process of a neural network

- Training pseudo code:

1. For each epoch:

1. For each data batch:

1. Perform a forward pass (prediction step)


2. Evaluate the error

3. Perform a backward pass (optimisation step)

4. Check termination condition

Introduction to model optimisation

- Derivative:
 - Measures the slope of the graph of a function at a given point
- Given a function $f(x)$, the derivative of the function is given by
 - $f'(x) = \lim_{\Delta x \rightarrow 0} \frac{f(x + \Delta x) - f(x)}{\Delta x} = \frac{\Delta y}{\Delta x}$
- To find a derivative of a function, the function must not contain any vertical slopes or breaks



Ref.2: Secant and tangent lines

Introduction to model optimisation

- Rules of Differentiation:

- Sum Rule: $\frac{d[f(x) + g(x)]}{dx} = \frac{df(x)}{dx} + \frac{dg(x)}{dx}$

- Difference Rule: $\frac{d[f(x) - g(x)]}{dx} = \frac{df(x)}{dx} - \frac{dg(x)}{dx}$

- Constant multiplication: $\frac{d[kf(x)]}{dx} = k \frac{df(x)}{dx}$

- Constant Rule : $\frac{d[C]}{dx} = 0$


- Product Rule: $\frac{d[f(x) \cdot g(x)]}{dx} = f(x) \frac{dg(x)}{dx} + g(x) \frac{df(x)}{dx}$

- The quotient Rule: $\frac{d[\frac{f(x)}{g(x)}]}{dx} = \frac{f(x) \frac{dg(x)}{dx} - g(x) \frac{df(x)}{dx}}{g^2(x)}$

Introduction to model optimisation

- Chain Rule of Derivation:
 - A composition function is:
 - $(f \circ g)(x) = f(g(x))$ and $(g \circ f)(x) = g(f(x))$
 - The chain rule is the derivative of a composition function
 - $\frac{d[f(g(x))]}{dx} = (f \circ g)'(x) = f'(g(x))g'(x)$


Introduction to model optimisation




Ref.3 Convex functions has a global minima/maxima and Non-Convex functions have multiple local minimas and a global minima

Introduction to model optimisation

- Function minima/maxima
 - On the ascending direction, the gradient is positive
 - On the descending direction, the gradient is negative
 - On maxima/minima the gradient is “0”
 - Goal of optimisation:
 - Find the function parameters that either minimise or maximise the function (parameters that leads to a $\frac{dy}{dx} = 0$)



Ref.2: Maxima of a function



Ref.2: Minima of a function

Introduction to model optimisation

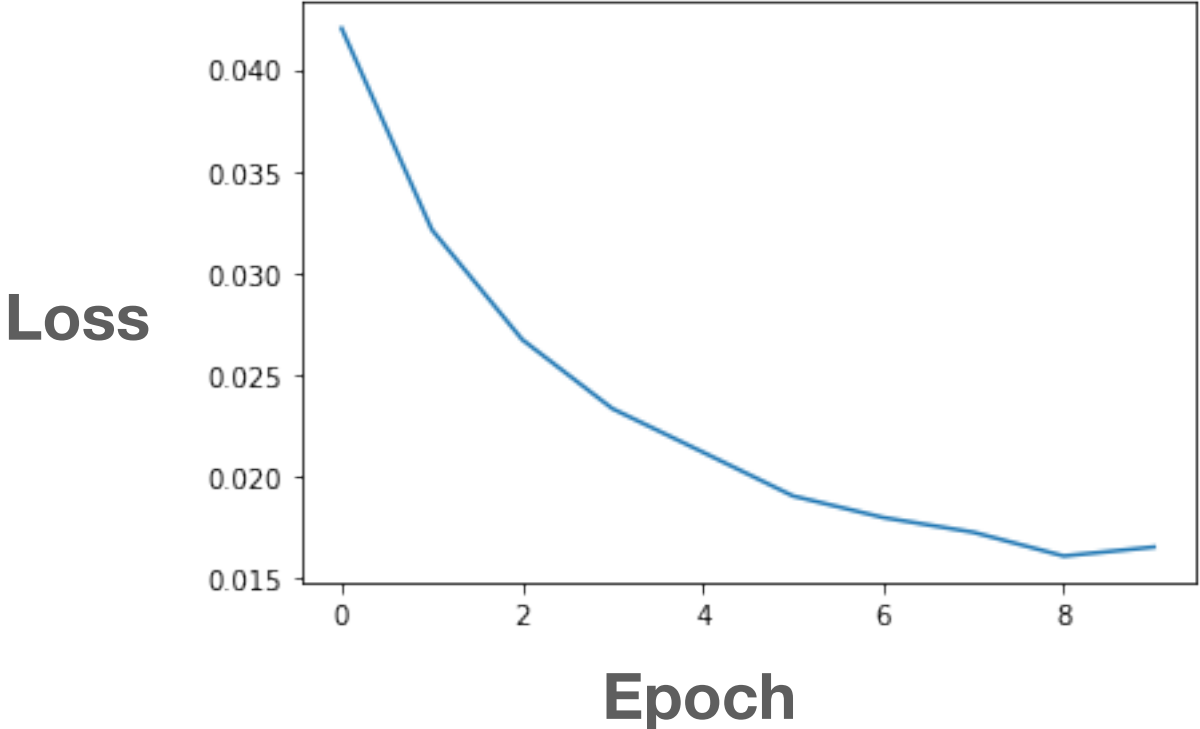
- Intuition on stepping towards a minima:
 - Take a function $f(\theta)$, a set of parameters θ and a loss function L
 - Parameter update equation:
 - $\theta_{new} = \theta_{old} - \alpha \frac{dL}{d\theta}$
 - $\frac{dL}{d\theta}$ is the loss derivative w.r.t parameters and α is the step size
 - So when the gradient is positive (ascending) we subtract the gradient so we move in the opposite direction
 - If the gradient is negative (descending) we keep moving on that direction (towards the minima)
 - Step size determines how big of a step we should take when calculating the next iteration parameters

Colab Exercise


Data Normalization

- What is Normalization:
 - Bring all features to a standard range while preserving their strength of representation
- Why normalise data:
 - Feature in the different ranges could poorly weigh outputs
 - The unnormalised data can lead to an unexpected behaviour in the gradients


Data Normalization




Train loss when inputs are 0-255




Train accuracy when inputs are 0-255



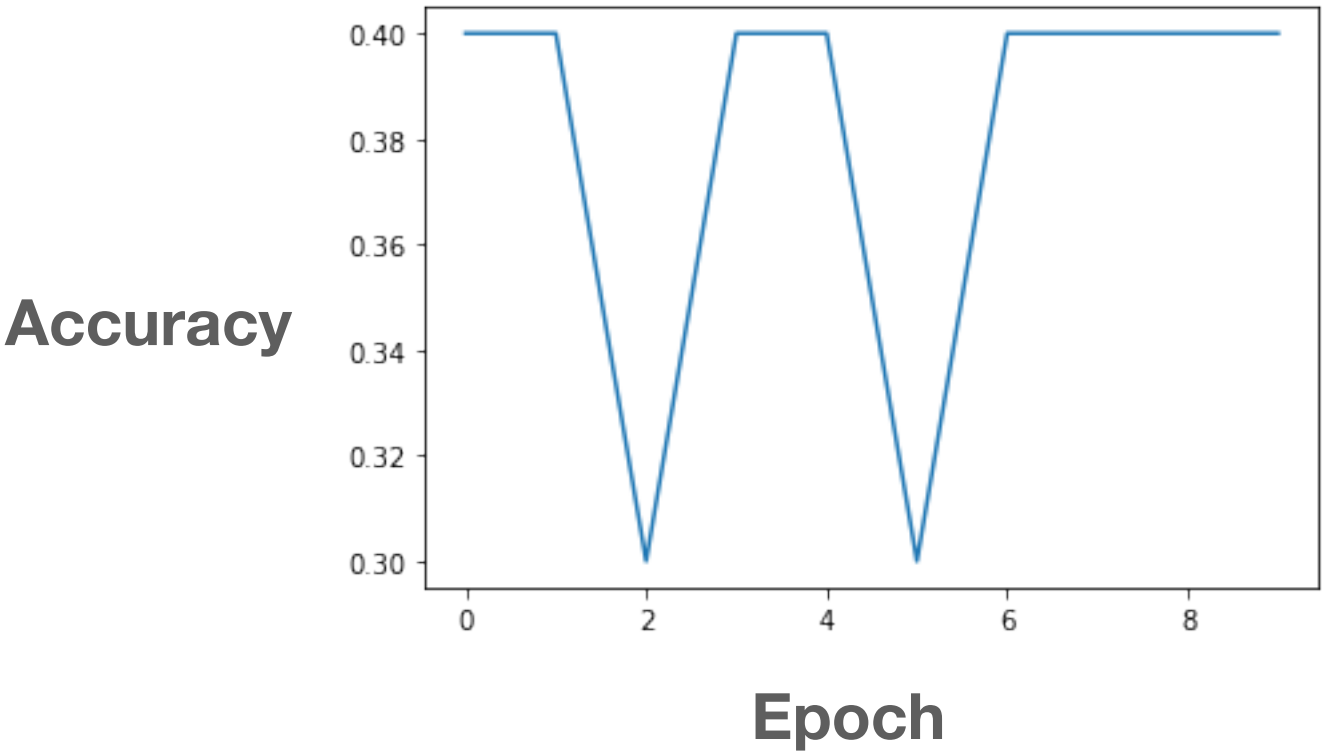
Test accuracy when inputs are 0-255



Train loss when inputs are 0-255*255=65025



Train accuracy when inputs are 0-255*255




Test accuracy when inputs are 0-255*255

Data Normalization

- What happens when data is not normalised:
 - Take gradient descent when there is no squashing function at output
 - When inputs are large the outputs will be large
 - Take a model $\hat{y} = Wx + b$, a sample loss: $L = \frac{1}{2}(\hat{y} - y)^2$ and its derivative $\frac{\partial L}{\partial W} = (\hat{y} - y)x$
 - So when x is large $\frac{\partial L}{\partial W}$ will be large

Data Normalization

- The parameter update is given by
$$w^{k+1} = w^k - \alpha \frac{\partial L}{\partial W}$$
- When the $\frac{\partial L}{\partial W}$ is very large, the update step would be very large
- This could overshoot the step and not converge towards a minima
- This is an exploding gradient issue
- If there is a squashing function like sigmoid, the gradients will vanish



Behaviour of the sigmoid derivative $\frac{\partial \sigma(x)}{\partial x}$ under different inputs

Data Normalization

- Data Normalization methods:

- Min/Max normalisation:

- Find the minimum and maximum of an input

- Normalise the input X using $X_{normalized} = \frac{X - X_{min}}{X_{max} - X_{min}}$

- Mean/Stdv normalisation (z-score method):

- Why?: When trained neural nets learn an underlying distribution of the input data

- By mean/stdv normalisation we make the distribution as simple as possible

- First subtract the mean from data (zero centering)

- Then divide by standard deviation (scale the data)

- $X_{normalized} = \frac{X - \mu}{\sigma}$

- Centering prevents vanishing gradients and scaling improves the convergence speed

Data Normalization

- Batch normalisation:

- Normalizes the inputs to a layer and applied on mini-batches

- Implemented in very deep neural networks

- Why?:

- Take two consecutive layers l and $l+1$.

- $l+1$ makes outputs based on the inputs from l

- In backprop $l+1$ is updated before l but based on the inputs from l and then the l is updated

- Now in the next iteration $l+1$ should model the distribution according to updated l

- This issue makes $l+1$ chase a moving target

- Therefore we need a standardisation of l layer inputs

- Batch normalisation is applied only during the training process

Data Normalization

- Batch normalisation equations:

- Input mini-batch: $B = [x_1 \dots x_m]$ where m is the batch size
- Output: $y_i = BN_{\gamma, \beta}(x_i)$ where γ, β are trainable parameters (backprop learning)
- γ, β trained to find the optimal distribution that minimises the error
- Computation

1. Find mean of the mini-batch: $\mu_B = \frac{1}{m} \sum_{i=1}^m x_i$

2. Find the variance of the mini-batch: $\sigma_B^2 = \frac{1}{m} \sum_{i=1}^m (x_i - \mu_B)^2$

3. Compute the normalised output: $\hat{x}_i = \frac{x_i - \mu_B}{\sqrt{\sigma_B^2 + \epsilon}}$ where ϵ is a regulariser to prevent division by “0”

4. Shift and scale the normalised output for optimal distribution: $y_i = \gamma \hat{x}_i + \beta = BN_{\gamma, \beta}(x_i)$

Reference

- Ref.1: <https://commons.wikimedia.org/wiki/File:MultiLayerPerceptron.svg>
- Ref.2: <https://medium.com/analytics-vidhya/concepts-of-differential-calculus-for-understanding-derivation-of-gradient-descent-in-linear-de59a17496a3>
- Ref.3: <https://slideplayer.com/slide/4916524/>
- Ref.4: <https://de.wikipedia.org/wiki/Datei:Sigmoid-function-2.svg>